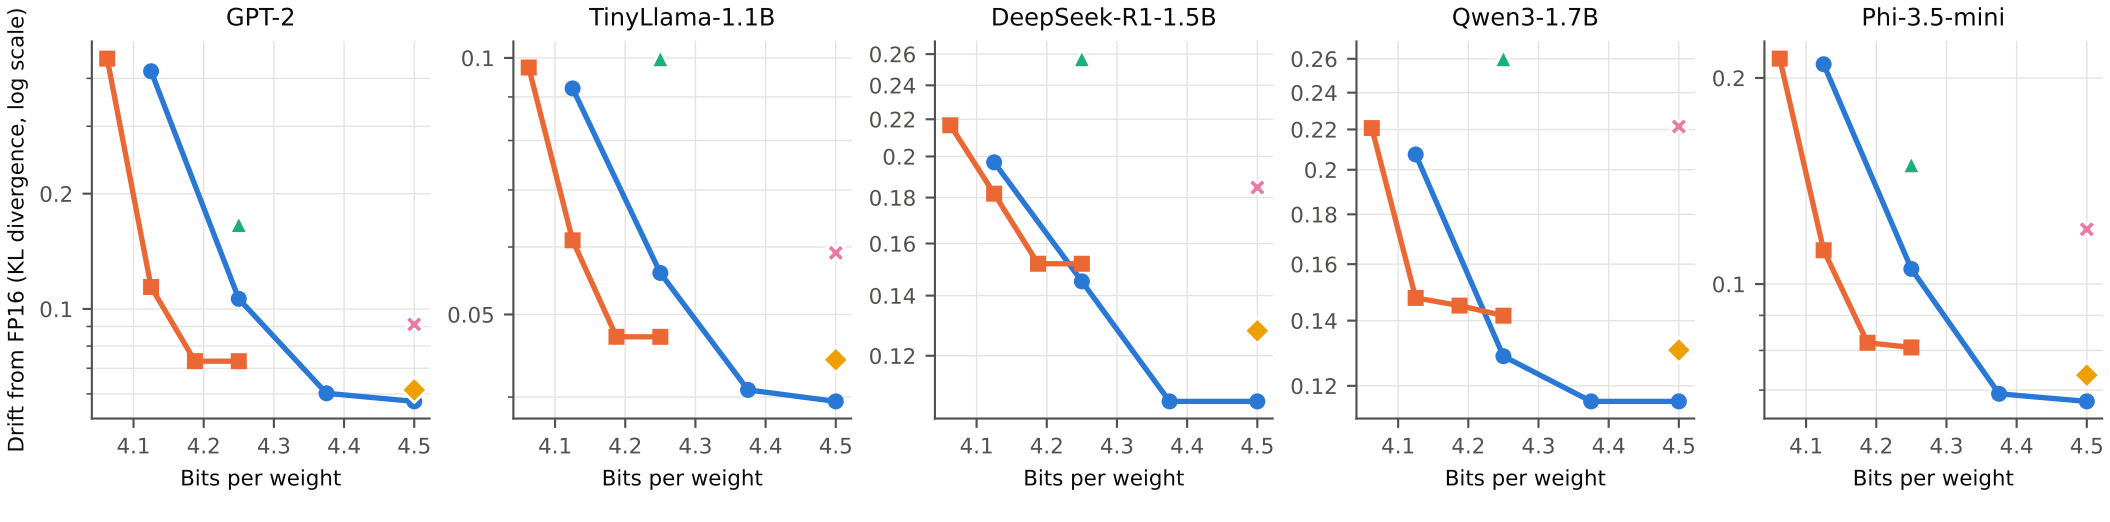
GPT-2
TinyLlama-1.1B
DeepSeek-R1-1.5B
Qwen3-1.7B
Phi-3.5-mini
Drift from FP16 (KL divergence, log scale)
Bits per weight
Scale sweep, block 16
Scale sweep, block 32
MXFP4 (published)
NVFP4 (published)
INT4, group 32